
SUMMARY

I am a curiosity driven, scientifically trained builder with over 10 years of experience in AI/ML, statistics, and graph algorithms. My academic work has received over 10,000 citations, and I've had the opportunity to contribute to production-grade systems in close collaboration with top-tier product teams. I am passionate about solving complex problems and building reliable, scalable tools across emerging areas like multi modal RAGs, intelligent search, and biomedical discovery.

EDUCATION**Stanford University***Ph.D. in Electrical Engineering · M.S. in Statistics*

Stanford, CA

2014 – 2020

Olin College of Engineering*B.S. in Electrical and Computer Engineering*

Needham, MA

2010 – 2014

EXPERIENCE**Nexa AI***Head of AI/ML*

Cupertino, CA

Feb 2025 – Present

- **RAG Productization:** Led a 4-member team to develop and launch an on-device RAG application in under 3 months, powered by continuous regression testing and weekly iteration to accelerate quality improvements.
- **Semantic Search Innovation:** Invented, prototyped, and shipped a semantic file-search engine with structured metadata, vector search, and @-search support, improving file recall quality in customer-facing RAG applications.

Apple*Machine Learning Engineer*

Cupertino, CA

Jan 2020 – Feb 2025

- **Offline Evaluation Service:** Spearheaded the development and rollout of an offline evaluation framework to continuously quantify feature impact on user experience prior to public releases, enabling faster improvement cycles. Owned multi-year roadmap strategy and stakeholder alignment. Scaled the initiative from a solo effort to a 5-person team, empowering members to independently drive new scopes and sustain the roadmap beyond my tenure.
- **Research Innovation:** Designed novel generative and retrieval-based frameworks to evaluate million-scale daily traffic in Apple Maps Search; shared methodologies at top-tier software engineering conferences (*ICSE*, *FSE*).
- **Infrastructure Modernization:** Revamped internal testing pipelines for query understanding and ranking, reducing release cycles from weekly to daily and improving launch stability for WWDC-featured products.
- **Technical Leadership:** Mentored team members in preparing and presenting their work at internal AI/ML conferences, fostering professional growth and cross-team visibility.

Stanford University*Graduate Research Assistant*

Stanford, CA

Sep 2014 – Feb 2020

- **Full-Stack Data Science:** Developed an interactive tool to visualize and perform power analysis on 30,000+ Gene Ontology terms—enabling large-scale association discovery with controlled false discovery rate.
- **Scalable Graph Learning:** Developed graph-based unsupervised learning pipelines for million-scale, high-dimensional datasets—resulting in publications in *Nature*, *Nature Methods*, *Cell*, and *NeurIPS*.

SELECTED PUBLICATIONS

1. Automatically Authoring Regression Tests for Machine-Learning-Based Systems. *ICSE*, 2021
2. Progenitor identification and SARS-CoV-2 infection in human distal lung organoids. *Nature*, 2020
3. Exploratory gene ontology analysis with interactive visualization. *Scientific Reports*, 2019
4. Visualization and analysis of sc-RNA-seq data by kernel-based similarity learning. *Nature Methods*, 2017

Full list shown on Google Scholar: <https://scholar.google.com/citations?user=2EasRdEAAAAJ&hl>